Integrated Results Report — Volatility-Reversal Finding

Volatility Reverses the Horizontal-Organization Advantage: How Environmental Turbulence Inverts the Information-Accessibility Prescription for Organizational Structure.

Working paper — simulation + preliminary archival pipeline.

Transparency note (2026-04-22)

An earlier draft of this work claimed cluster-based hybrid dominance under environmental volatility. That claim was wrong. The earlier dynamic-experiment code delegated the hybrid trajectory to `v2.run_hybrid`, which does not apply landscape shocks. When shocks were correctly applied to all three structures, the hybrid advantage disappeared. We re-ran the experiments, revised the paper, and document the correction openly.

The corrected finding is reported below.

Executive summary

Four findings from **36,900 ABM runs across nine pre-registered batteries**:

1. **Static baseline confirms classical contingency.** In static NK landscapes, vertical wins at low τ, horizontal wins at high τ, and the crossover τ*(K) decreases in task complexity. H1–H4 supported.
2. **Headline — volatility reverses the horizontal advantage.** Introducing environmental shocks monotonically shifts the V-vs-H performance gap toward vertical. At one shock per 100 steps (near-static), H beats V by 0.029 (95% bootstrap CI [−0.036, −0.023]). At one shock per 3 steps (high-vol), V beats H by 0.081 ([+0.070, +0.092]). Slope on log(ω): +0.030 (95% CI [+0.026, +0.033]); permutation $p < 0.002$; Spearman $\rho = +1.00$ across six levels.
3. **The headline survives aggressive robustness checks.** The V-dominance pattern replicates across three independent seed banks ($\rho \geq +0.88$), four alternative shock-model specifications (periodic/Poisson/drift/correlated, $\rho \geq +0.94$), four agent-heterogeneity regimes from homogeneous to extreme-star-distribution (all $\rho = +1.00$), and a multi-level vertical structure with local-authority mid-managers (Burns-Stalker mechanistic form) that rules out a "dictatorship-V" artifact.
4. **Hybrid is mid-pack.** Performs between V and H in both regimes. Modularity does not uniformly dominate.

---

1. Static baseline (H1–H4)

Results from `results.csv` (v1 aggregation-rule), `results_v2.csv` (v2 independent-exploration with fragmentation cost), and `results_robustness.csv` (7,020-run sensitivity battery).

v2 (headline static model)

- Clean crossovers: τ*(K=2) = 0.16, τ*(K=5) = 0.13, τ*(K=10) = 0.13.
- Fine-grid refinement: τ*(K=2) = 0.157, τ*(K=5) = 0.121, τ*(K=10) = 0.119.
- H − V gap peaks at +0.055 for K = 10, τ = 0.3.

- H4 ($\partial\tau^*/\partial K < 0$) quantitatively confirmed.

Figure: `results/figures/fig_delta_HV_v2.png`

Robustness sweeps

- Fragmentation cost $\in \{0.10, 0.20, 0.30\}$: shifts low-τ endpoint but crossover structure persists.
- n_agents $\in \{15, 30, 60\}$: qualitative patterns invariant.
- Fine-τ grid (step 0.025): refined τ^* estimates monotone-decreasing in K.

---

2. Volatility reversal (H5 — headline)

Fixed shock frequency (1,620 runs, SHOCK_EVERY = 10)

Results in `results_dynamic.csv`. V uniformly outperforms H across (τ, K). Hybrid between.

K	τ	V	H	Hy	V − H
2	0.1	0.682	0.515	0.589	+0.167
2	0.9	0.676	0.702	0.693	−0.026
5	0.1	0.700	0.539	0.594	+0.161
5	0.5	0.690	0.634	0.645	+0.056
10	0.1	0.659	0.525	0.553	+0.134
10	0.5	0.685	0.603	0.676	+0.082

Figure: `results/figures/fig_perf_curves_dynamic.png`

Volatility gradient (4,860 runs, SHOCK_EVERY $\in \{3, 5, 10, 20, 40, 100\}$)

Results in `results_volatility.csv`. Cleanest finding of the paper:

SHOCK_EVERY	Volatility ω	V − H mean	SEM
100	0.010	**−0.029**	0.005
40	0.025	+0.021	0.010
20	0.050	+0.046	0.014
10	0.100	+0.054	0.015
5	0.200	+0.071	0.023
3	0.333	+0.081	0.023

Spearman ρ(ω, V − H) = +1.00 across the six levels.

Crossover between shock/100 (H wins by 0.029) and shock/40 (V wins by 0.021). Above the crossover, V dominates uniformly.

Figures:
- `results/figures/fig_VH_gap_by_volatility.png` — headline dose-response
- `results/figures/fig_vol_by_tau.png` — by-K and τ breakdown
- `results/figures/fig_structure_fitness_vol.png` — all three structures' absolute performance vs volatility

3. Robustness battery (2026-04-23 extension)

Four orthogonal critiques of the headline were addressed by running four additional experiments. All four confirm H5.

3.1 Statistical robustness — 3 independent seed banks (4,860 runs)

Results in `results_replication.csv`. Banks A, B, C use disjoint seed ranges.

Seed bank	V−H at shock/100	V−H at shock/3	Slope on log(ω)	Spearman ρ
A	−0.024	+0.089	+0.0292	+1.000
B	−0.032	+0.089	+0.0316	+0.943
C	−0.040	+0.082	+0.0323	+0.886

All three banks yield monotone positive dose-response. **Figure:** `results/figures/fig_replication_3banks.png`.

3.2 Shock-model specificity — 4 variants (6,480 runs)

Results in `results_shocks.csv`. The baseline 25%-locus-reshuffle is only one operationalization; we test four.

Shock model	Min-intensity V−H	Max-intensity V−H	Spearman ρ	Perm p
periodic (baseline)	−0.034	+0.083	+0.943	0.0087
poisson (random timing)	−0.023	+0.077	+0.943	0.0087
drift (continuous small perturbations)	−0.032	+0.039	+1.000	0.0012
correlated (block shifts)	−0.026	+0.057	+1.000	0.0012

All four give the same qualitative pattern. H5 is not specific to the periodic reshuffle. **Figure:** `results/figures/fig_shocks_comparison.png`.

3.3 Multi-level vertical — Burns-Stalker mechanistic form (4,860 runs)

Results in `results_multilevel.csv`. The baseline V is effectively a CEO dictatorship; ML-V adds 5 mid-managers with local authority and a CEO coordination step every 3 time units.

shock_every	V	ML-V	H	ML−H	V−H
100	0.700	0.703	0.733	−0.030	−0.033
40	0.680	0.678	0.651	+0.027	+0.030
20	0.677	0.676	0.645	+0.031	+0.032
10	0.682	0.673	0.624	+0.049	+0.059
5	0.678	0.681	0.611	+0.070	+0.066
3	0.690	0.698	0.606	+0.092	+0.084

ML-V tracks V closely at every volatility level. The V-dominance is a property of hierarchical structure, not of centralization extremity. **Figure:** `results/figures/fig_multilevel_V.png`.

3.4 Heterogeneous agents — 4 talent regimes (720 runs)

Results in `results_heterogeneous.csv`. Real organizations have large variance in individual competence; the critique was that "V wins under volatility" might really be "talent-concentration wins."

Heterogeneity regime	V−H at static	V−H at high vol	Spearman ρ
homogeneous (all agents share τ)	−0.036	+0.083	+1.000
mild (τ_i ~ Normal(τ, 0.10))	−0.040	+0.062	+1.000
skewed (20% stars at 2τ, 80% at 0.5τ)	−0.050	+0.115	+1.000
extreme (10% at τ=1.0, 90% at 0.3τ)	−0.031	+0.158	+1.000

All four regimes yield perfect ρ = +1.00. The **extreme** regime actually amplifies the V-advantage at high volatility (+0.158 vs homogeneous +0.083), but the monotone dose-response survives in every distribution. H5 is not a talent-concentration artifact. **Figure:** `results/figures/fig_heterogeneous_agents.png`.

3.5 What the robustness battery establishes

Under nine orthogonal perturbations — three seed banks, four shock models, a structurally-different vertical form, and four talent distributions — H5 never reverses, never flattens, and never falls below Spearman ρ = +0.88. This is about as strong as ABM evidence gets.

---

4. Archival pipeline — preliminary

N = 111 firm-years, 47 firms, from S&P 500 (2021–2025). Data from SEC EDGAR 10-K filings + yfinance.

Test	Estimate	p-value	Honest read
H1 (flat coef, low-τ n=37)	+0.322	0.032	wrong sign; subsample too small
H2 (flat coef, high-τ n=37)	+0.060	0.303	right direction, not significant
H3 (flat × τ, n=111)	−0.012	0.961	null
H4 (flat × τ × K, n=111)	+92.21	<0.001	artifact from K-variance collapse

H5 cannot be tested on the current panel — we lack an industry-level or firm-level volatility proxy.

Pipeline is validated; substantive inference requires the sample expansion listed in the manuscript (Russell 3000 + DEF 14A + Compustat + volatility proxy).

---

5. Implications

For contingency theory

Burns and Stalker's (1961) claim that organic structures fit turbulent environments is not supported by our NK operationalization. The direction is the opposite: turbulence favors vertical's concentrated, single-path decision-making over horizontal's distributed parallel search.

For the self-managing-organization literature

Lee and Edmondson (2017) concluded qualitatively that radical-flat firms endure under "specific configurations." Our simulation makes that precise: those configurations require both (a) high τ and (b) low environmental volatility.

Morning Star (tomato paste, stable production) and Buurtzorg (nursing care, slow-changing) meet both; Zappos, Oticon, Valve (consumer markets, fast change) do not.

For AI-era organizational design

The popular "AI + volatility → flat" argument has two steps. The first (AI raises τ) is supported by our static ABM. The second (volatility favors flat) is **contradicted** by our dynamic ABM. Any claim that AI will flatten organizations must separately argue that AI raises τ faster than it raises ω.

For empirical research design

Future studies comparing flat and hierarchical firm performance should separately measure industry-level environmental volatility (margin variance, regulatory-change frequency). Our simulation predicts that flat-firm advantages found at high τ in stable industries will not generalize to high-volatility industries.

---

6. Reproducibility manifest

Asset	Location	How to regenerate
Static ABM v1	`results/results.csv`	`python abm/experiment.py`
Static ABM v2	`results/results_v2.csv`	`python abm/experiment_v2.py`
Parameter robustness	`results/results_robustness.csv`	`python abm/robustness.py`
Dynamic ABM (fixed shock)	`results/results_dynamic.csv`	`python abm/experiment_dynamic.py`
Volatility gradient (headline)	`results/results_volatility.csv`	`python abm/experiment_volatility.py`
3-bank replication	`results/results_replication.csv`	`python abm/experiment_replication.py`
Alternative shock models	`results/results_shocks.csv`	`python abm/experiment_shocks.py`
Multi-level V	`results/results_multilevel.csv`	`python abm/experiment_multilevel.py`
Heterogeneous agents	`results/results_heterogeneous.csv`	`python abm/experiment_heterogeneous.py`
Bootstrap 95% CIs	`results/bootstrap_summary.json`	`python abm/bootstrap_analysis.py`
All figures	`results/figures/*.png`	`python abm/visualize_robustness_all.py` (also calls the three baseline visualize scripts)
Archival panel	`empirical/real_data/real_panel.csv`	EDGAR + yfinance + parse_fix.py + merge_panel.py
Regression summary	`empirical/real_data/regression_real_summary.csv`	`python empirical/real_data/regression_real.py`
Manuscript PDF	`submission/manuscript.pdf`	`python abm/final_assembly.py` (fills placeholders, regenerates figures, rebuilds PDF)

All seeds fixed. ~60–90 min total replication on a laptop (~half that if the four extensions run in parallel).